

# STOR 455 Homework #5

40 points - Due Wednesday 3/20 at 5:00pm

**Directions:** For parts 7 and 10 you should work together, but these parts must be **submitted individually** by each group member. For parts 8 and 9, you must have only **one submission per group**. There will be separate places on Gradescope to submit the individual vs group work.

**Situation:** Can we predict the selling price of a house in Ames, Iowa based on recorded features of the house? That is your task for this assignment. Each team will get a dataset with information on forty potential predictors and the selling price (in \$1,000's) for a sample of homes. The data sets for your group are AmesTrain??.csv and AmesTest??.csv (where ?? corresponds to your group number) A separate file identifies the variables in the Ames Housing data and explains some of the coding.

**Part 7. Cross-validation:** In some situations, a model might fit the peculiarities of a specific sample of data well, but not reflect structure that is really present in the population. A good test for how your model might work on “real” house prices can be simulated by seeing how well your fitted model does at predicting prices that were NOT in your original sample. This is why we reserved an additional 200 cases as a holdout sample in AmesTest??.csv. Use the group number and AmesTest??.csv corresponding to your group number for homework #3. Import your holdout test data and

- Compute the predicted Price for each of the cases in the holdout test sample, using your model resulting from the initial fit and residual analysis in parts 1 through 3 of Homework #3.

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.4
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.4.4      v tibble    3.2.1
## v lubridate  1.9.3      v tidyr     1.3.0
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

AmesData <- read.csv('AmesTrain10.csv')

AmesData1 <- select(AmesData, is.numeric)

## Warning: Use of bare predicate functions was deprecated in tidyselect 1.1.0.
## i Please use wrap predicates in `where()` instead.
##   # Was:
##   data %>% select(is.numeric)
##
##   # Now:
##   data %>% select(where(is.numeric))
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```

```
holdout <- read.csv('AmesTest10.csv')

reducedmodel1 <- lm(formula = Price ~ Quality + GroundSF + BasementFinSF + BasementSF +
  LotArea + YearBuilt + GarageSF + YearRemodel + Bedroom +
  LotFrontage + FullBath + Condition, data = AmesData1)

test_data <- predict(reducedmodel1, newdata = holdout)
```

- Compute the residuals for the 200 holdout cases.

```
test_resid <- holdout$Price - test_data
```

- Compute the mean and standard deviation of these residuals. Are they close to what you expect from the training model?

```
mean(test_resid)
```

```
## [1] -1.631249
```

```
sd(test_resid)
```


```
## [1] 28.98615
```

- Construct a plot of the residuals to determine if they are normally distributed. Is this plot what you expect to see considering the training model?

```
qqnorm(test_resid)
```

```
qqline(test_resid)
```

## Normal Q-Q Plot



- Are any holdout cases especially poorly predicted by the training model? If so, identify by the row number(s) in the holdout data. Why might these cases be poorly predicted?

```
sort(abs(test_resid), decreasing = TRUE)[1:5]
```

```
##      118      48      57      175      186
```

```
## 179.10581 120.69174 114.15403 98.29246 90.12243
```

- Compute the correlation between the predicted values and actual prices for the holdout sample. This is known as the cross-validation correlation. We don't expect the training model to do better at predicting values different from those that were used to build it (as reflected in the original  $R^2$ ), but an effective model shouldn't do a lot worse at predicting the holdout values. Square the cross-validation correlation to get an  $R^2$  value and subtract it from the original multiple  $R^2$  of the training sample. This is known as the shrinkage. We won't have specific rules about how little the shrinkage should be, but give an opinion on whether the shrinkage looks OK to you or too large in your situation.

```
summary(reducedmodel1)$r.squared - cor(holdout$Price,test_data)^2
```

```
## [1] 0.0331965
```

**Part 8. Find a “fancy model”:** Again using AmesTrain`??`.csv, where `??` corresponds to your new group number in homework #5, to build a regression model to predict Price. In addition to the quantitative predictors, you may now consider models with

```
NewAmesTrain <- read.csv('AmesTest10.csv', stringsAsFactors = TRUE)
NewAmesTrain <- na.omit(NewAmesTrain)
NewAmesTrain <- NewAmesTrain[-1]
age_built <- 2010 - NewAmesTrain$YearBuilt
age_remodel <- 2010 - NewAmesTrain$YearRemodel

full <- lm(Price ~ ., data=NewAmesTrain)
#summary(full)
```

- Categorical variables - Just put these in the model and let R take care of making the indicator predictors (and picking one category to leave out). Use `factor()` to treat a numeric variable as categorical. You'll see the coefficients for each indicator when you look at the `summary()` and they will be grouped together in the ANOVA. Be careful, since adding a single categorical variable with a lot of categories might actually be adding a lot of new indicator terms.

```
NewAmesTrain_wfactor <- NewAmesTrain
NewAmesTrain_wfactor$Quality <- factor(NewAmesTrain_wfactor$Quality)
NewAmesTrain_wfactor$Condition <- factor(NewAmesTrain_wfactor$Condition)

full_wfactor <- lm(Price ~ ., data=NewAmesTrain_wfactor)
summary(full_wfactor)
#plot(Price ~ ., data=NewAmesTrain_wfactor)
```

- Transformations of predictors - You can include functions of quantitative predictors. Probably best to use the `I()` notation so you don't need to create new columns when you run the predictions for the test data.
- Transformations of the response - You might address curvature or skewness in residual plots by transforming the response prices with a function like `log(Price)`, `sqrt(Price)`, `Price^2`, etc.. These should generally not need the `I()` notation to make these adjustments. IMPORTANT: If you transform Price, be sure to reverse the transformation when making final predictions!
- Combinations of variables - This might include interactions or other combinations. You do not need the `I()` notation when making an interaction using a categorical predictor (e.g. `GroundSF*CentralAir`).

Keep general track of the approaches you try and explain what guides your decisions as you select a new set of predictors (but again you don't need to give full details of every model you consider). Along the way you should consider some residual analysis.

Notes/Tips:

- WARNING: When using a categorical predictor with multiple categories in `regsubsets()`, R will create indicators and treat them as separate predictors when deciding which to put into a model. So you might

get a model with quantitative predictors like LotArea and GroundSF along with specific indicators like GarageQTA and HouseStyle1Story. This may not be very useful, since we should generally use all indicators for a categorical predictor if we include one in the model. On the other hand, when using the `step()` function, R will generally keep the multiple indicators for different categories of the same variable together as a unit.

- In some cases the indicators created for different categorical variables will have identical values. For example, if you include both GarageC and GarageQ in a model, R will produce values for each of the indicators. The indicators for GarageQNone and GarageCNone (equal to one only for houses that don't have a garage) will be identical. This may be handled differently in R depending on the procedure. `regsubsets()` may give a “warning” about variables being linearly dependent. You can still use the results, just be aware that some variables are completely dependent. `lm()` might give output with coefficients (and tests) of some predictors listed as NA. This is not a problem, R is just automatically deleting one of the redundant variables. If you are predicting for a house with no garage you might have a coefficient to use for GarageQNone but then you don't need to worry about having one for GarageCNone.
- If your residual analysis from homework #3 or an early model here suggest you might want to do a transformation for the response variable (Price), do so *before* fitting a lot more models. No sense fine tuning a set of predictors for Price, then deciding you should be predicting  $\log(\text{Price})$  or  $\text{Price}^2$ . So make that decision fairly early, but don't get too picky and expect to get perfect plot of residuals versus fits or an exact normal quantile plot.
- Similarly, if you decide that some data cases should be dropped from the training set, don't wait until late in the process to do so. For example, if you spot a *very* large residual you should look at the characteristics for that house to see if it should be deleted. Don't forget about the value of simple plots (like a scatterplot of Price vs. LotArea) for helping to see what is going on and recognize extreme cases. Be sure to document any adjustments you make in the final report.
- Comparing  $C_p$  from different predictor pools - While Mallows's  $C_p$  is a useful tool for comparing models from the same pool of predictors. You should not use it to compare models based on different predictor pools. For example, if you add a bunch of categorical variables to all the quantitative predictors from homework #3 to make a new “full” model, then find  $C_p$  from a model that you fit in homework #3, it will be worse than it was before. If you look at the formula for calculating  $C_p$ , you will see that all that has changed is MSE for the full model after adding the new batch of predictors.
- I should be able to follow the steps you use when selecting a model. I certainly don't need to see every bit of output, but it might help to include more of the R commands you use. For example, saying you used backward elimination is not very helpful when I don't know what you start with for the full model or pool of predictors (e.g. did you include Condition and Quality as numeric predictors? or did you decide to eliminate one of GroundSF, FirstSF, or SecondSF due to redundancy?). The easiest way to convey this in many cases is to show the R command you used. It is fine to abbreviate the output (for example, delete many steps in a stepwise procedure using `trace=FALSE`), but it would be helpful if you identified the parts you do include. For example, a sentence like “After 12 steps of the stepwise procedure, we have the output below for the fitted model.” Similarly, I don't need to see 600 residuals, using `head` and `sort` can show the important ones.
- Once you have settled on a response, made adjustments to the data (if needed), and chosen a set of predictors, be sure to include the `summary()` for your “fancy” model at this stage.

```
DoublenewNewAmesTrain <- NewAmesTrain_wfactor %>%
  mutate(Price = log(Price),
         LotArea = log(LotArea))

first_model <- lm(log(Price) ~ . + I(age_built^2) + I(GarageCars^2) + I(GarageSF^2) + TotalRooms * Bedroom
summary(first_model)

##
## Call:
```

```
## lm(formula = log(Price) ~ . + I(age_built^2) + I(GarageCars^2) +
##      I(GarageSF^2) + TotalRooms * Bedroom + LotArea * LotFrontage +
##      GarageSF * GarageCars, data = DoublenewNewAmesTrain)
##
## Residuals:
##      Min        1Q      Median        3Q       Max
## -0.037757 -0.009151  0.000000  0.008070  0.056877
##
## Coefficients: (7 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2.549e-01  8.815e-01   0.289 0.773003
## LotFrontage   -1.189e-03  1.201e-03  -0.990 0.324216
## LotArea       1.609e-02  5.878e-03   2.737 0.007257 **
## LotConfigCulDSac 1.185e-03  8.151e-03   0.145 0.884657
## LotConfigFR2   -1.565e-02  1.044e-02  -1.498 0.136932
## LotConfigFR3   -2.926e-02  2.940e-02  -0.995 0.321838
## LotConfigInside  8.673e-04  4.667e-03   0.186 0.852929
## HouseStyle1Story  5.682e-03  9.855e-03   0.577 0.565436
## HouseStyle2.5Unf -1.600e-02  2.215e-02  -0.723 0.471502
## HouseStyle2Story -1.219e-02  8.393e-03  -1.453 0.149182
## HouseStyleSFoyer -1.448e-03  1.246e-02  -0.116 0.907722
## HouseStyleSLvl  1.322e-03  1.214e-02   0.109 0.913512
## Quality4       2.856e-02  1.679e-02   1.701 0.091848 .
## Quality5       5.478e-02  1.680e-02   3.261 0.001488 **
## Quality6       6.157e-02  1.716e-02   3.588 0.000503 ***
## Quality7       7.508e-02  1.791e-02   4.191 5.70e-05 ***
## Quality8       8.514e-02  1.926e-02   4.421 2.35e-05 ***
## Quality9       1.304e-01  2.503e-02   5.211 9.09e-07 ***
## Condition4     1.619e-01  3.331e-02   4.862 3.98e-06 ***
## Condition5     1.724e-01  3.265e-02   5.281 6.72e-07 ***
## Condition6     1.795e-01  3.225e-02   5.567 1.91e-07 ***
## Condition7     1.908e-01  3.279e-02   5.820 6.11e-08 ***
## Condition8     1.992e-01  3.324e-02   5.994 2.75e-08 ***
## Condition9     2.333e-01  3.759e-02   6.205 1.03e-08 ***
## YearBuilt      4.640e-04  5.139e-04   0.903 0.368554
## YearRemodel    -2.584e-05  1.691e-04  -0.153 0.878799
## ExteriorQFa    -3.655e-03  4.731e-02  -0.077 0.938555
## ExteriorQGd     2.803e-02  1.883e-02   1.489 0.139498
## ExteriorQTA     3.429e-02  1.987e-02   1.725 0.087303 .
## ExteriorCFa    -6.090e-03  3.290e-02  -0.185 0.853521
## ExteriorCGd    -3.451e-02  2.639e-02  -1.308 0.193737
## ExteriorCTA    -1.788e-02  2.612e-02  -0.685 0.495029
## FoundationCBlock  5.478e-03  9.237e-03   0.593 0.554422
## FoundationPConc  1.517e-02  9.726e-03   1.559 0.121815
## FoundationSlab   4.545e-02  2.680e-02   1.696 0.092755 .
## FoundationStone  5.671e-02  2.557e-02   2.218 0.028635 *
## FoundationWood  -3.488e-02  2.698e-02  -1.292 0.198978
## BasementHtEx    7.754e-03  1.131e-02   0.686 0.494455
## BasementHtFa    1.041e-02  1.507e-02   0.691 0.491268
## BasementHtGd    6.147e-03  6.808e-03   0.903 0.368536
## BasementHtNone  -8.456e-03  2.771e-02  -0.305 0.760830
## BasementHtTA      NA         NA         NA         NA
## BasementCFa     -7.417e-03  1.285e-02  -0.577 0.565105
## BasementCGd     -5.372e-04  6.765e-03  -0.079 0.936851
## BasementCNone      NA         NA         NA         NA
```

## BasementCTA	NA	NA	NA	NA
## BasementFinALQ	8.061e-04	8.288e-03	0.097	0.922700
## BasementFinBLQ	-1.177e-03	8.922e-03	-0.132	0.895325
## BasementFinGLQ	6.090e-03	7.536e-03	0.808	0.420740
## BasementFinLwQ	-8.512e-05	1.024e-02	-0.008	0.993381
## BasementFinNone	NA	NA	NA	NA
## BasementFinRec	-3.509e-03	1.047e-02	-0.335	0.738189
## BasementFinUnf	NA	NA	NA	NA
## BasementFinSF	7.867e-06	1.530e-05	0.514	0.608187
## BasementUnFinSF	-3.323e-06	1.565e-05	-0.212	0.832293
## BasementSF	4.842e-05	2.231e-05	2.170	0.032190 *
## HeatingGasW	3.306e-03	1.512e-02	0.219	0.827378
## HeatingQCFA	-5.978e-03	1.341e-02	-0.446	0.656746
## HeatingQCGd	3.582e-03	5.696e-03	0.629	0.530717
## HeatingQCTA	-7.864e-04	6.049e-03	-0.130	0.896799
## CentralAirY	9.262e-04	1.207e-02	0.077	0.938995
## FirstSF	7.436e-05	5.611e-05	1.325	0.187946
## SecondSF	1.130e-04	5.833e-05	1.938	0.055250 .
## GroundSF	-4.273e-05	5.421e-05	-0.788	0.432299
## BasementFBath	-7.958e-05	5.255e-03	-0.015	0.987946
## BasementHBath	-2.250e-03	1.304e-02	-0.173	0.863333
## FullBath	1.524e-03	6.154e-03	0.248	0.804863
## HalfBath	-5.233e-03	6.119e-03	-0.855	0.394329
## Bedroom	-7.439e-03	1.112e-02	-0.669	0.504933
## KitchenQFA	-7.219e-02	2.499e-02	-2.889	0.004677 **
## KitchenQGd	-3.507e-03	9.821e-03	-0.357	0.721743
## KitchenQTA	-1.161e-02	1.109e-02	-1.047	0.297615
## TotalRooms	4.379e-03	5.170e-03	0.847	0.398837
## Fireplaces	5.128e-03	3.725e-03	1.377	0.171477
## GarageTypeAttchd	-1.615e-02	3.011e-02	-0.537	0.592707
## GarageTypeBasment	2.585e-04	3.488e-02	0.007	0.994100
## GarageTypeBuiltIn	-2.466e-02	3.154e-02	-0.782	0.436131
## GarageTypeCarPort	-6.584e-02	3.503e-02	-1.879	0.062896 .
## GarageTypeDetchd	-7.214e-03	2.858e-02	-0.252	0.801211
## GarageTypeNone	-5.605e-02	5.011e-02	-1.119	0.265760
## GarageCars	5.521e-03	1.993e-02	0.277	0.782308
## GarageSF	-1.228e-05	7.075e-05	-0.174	0.862473
## GarageQGd	-6.876e-03	1.972e-02	-0.349	0.728045
## GarageQNone	NA	NA	NA	NA
## GarageQTA	-2.032e-02	1.276e-02	-1.593	0.114182
## GarageCFA	3.552e-02	3.176e-02	1.118	0.265914
## GarageCGd	-4.147e-03	3.806e-02	-0.109	0.913443
## GarageCNone	NA	NA	NA	NA
## GarageCTA	2.467e-02	2.843e-02	0.868	0.387468
## WoodDeckSF	-1.288e-06	1.650e-05	-0.078	0.937917
## OpenPorchSF	7.510e-07	2.967e-05	0.025	0.979849
## EnclosedPorchSF	-5.379e-06	3.137e-05	-0.171	0.864172
## ScreenPorchSF	-1.704e-05	4.245e-05	-0.401	0.688952
## I(age_built^2)	-2.133e-06	4.228e-06	-0.504	0.615020
## I(GarageCars^2)	-1.267e-03	1.018e-02	-0.124	0.901231
## I(GarageSF^2)	1.638e-08	1.353e-07	0.121	0.903868
## Bedroom:TotalRooms	9.633e-05	1.513e-03	0.064	0.949345
## LotFrontage:LotArea	1.230e-04	1.278e-04	0.963	0.337778
## GarageCars:GarageSF	-1.532e-06	6.449e-05	-0.024	0.981094
## ---				

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01959 on 108 degrees of freedom
## Multiple R-squared:  0.9602, Adjusted R-squared:  0.9267
## F-statistic: 28.64 on 91 and 108 DF,  p-value: < 2.2e-16
```

#Firstly, we opted to transform the price variable by taking its logarithm. This decision was based on its ability to mitigate the influence of extreme values, thus enhancing the fit of our model, particularly with respect to its curvature.

#Similarly, we applied a logarithmic transformation to the LotArea variable. Our rationale behind this choice stems from observing a curving trend within the data, which logarithmic adjustment tends to alleviate effectively.

#For variables such as age built and garage cars, which exhibited polynomial trends according to our visual inspection of scatter plots, we decided to employ polynomial transformations. This method allowed us to capture the non-linear relationships more accurately, thereby improving the model's performance.

#Additionally, recognizing the high correlation between total rooms and total bedrooms, we deemed it essential to include an interaction term between these two variables. By doing so, we aimed to capture any nuanced interplay between these features, thus enriching the model's predictive capacity.

#Similarly, we identified a significant correlation between lot area and lot frontage, prompting us to incorporate an interaction term for these variables as well. This approach enables us to account for their mutual influence more effectively, enhancing the model's robustness.

#Finally, given the strong correlation between garage cars and garage square footage, we decided to handle them in a similar manner, introducing an interaction term to capture their combined effect accurately.

```
MSE2 <- (summary(first_model)$sigma)^2

AmesModNone2 <- lm(Price ~ 1, NewAmesTrain_wfactor)

#simple_model<-step(first_model, scale = MSE2)
#step(AmesModNone2, scope = list(upper=first_model), scale=MSE2, direction='forward')
simple_model <- step(AmesModNone2, scale = MSE2, scope = list(upper = first_model), trace = FALSE)

summary(simple_model)
```

```
##
## Call:
## lm(formula = Price ~ Quality + GroundSF + Condition + YearBuilt +
##     BasementFinSF + ExteriorQ + HouseStyle + GarageSF + Foundation +
##     LotArea + LotConfig + GarageType + Bedroom + TotalRooms +
##     ExteriorC + BasementHBath + I(GarageSF^2) + GarageC + HeatingQC +
##     BasementSF + BasementC + BasementHt + KitchenQ + GarageQ +
##     ScreenPorchSF + Fireplaces + WoodDeckSF + FullBath + HalfBath +
##     SecondSF + FirstSF + BasementFin + BasementUnFinSF + BasementFBath +
##     YearRemodel + Heating + EnclosedPorchSF + LotFrontage + I(age_built^2) +
##     GarageCars + I(GarageCars^2) + CentralAir + OpenPorchSF +
##     Bedroom:TotalRooms + LotArea:LotFrontage + GarageSF:GarageCars,
##     data = NewAmesTrain_wfactor)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -46.069  -9.990   0.000   8.393  74.343
##
## Coefficients: (7 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -1.767e+03  9.761e+02  -1.810 0.073036 .
```

## Quality4	-3.349e+00	1.854e+01	-0.181	0.856997
## Quality5	-2.475e+00	1.851e+01	-0.134	0.893861
## Quality6	7.226e+00	1.892e+01	0.382	0.703224
## Quality7	1.704e+01	1.977e+01	0.862	0.390433
## Quality8	3.056e+01	2.102e+01	1.454	0.148874
## Quality9	5.414e+01	2.717e+01	1.993	0.048805 *
## GroundSF	-1.442e-01	5.930e-02	-2.431	0.016699 *
## Condition4	1.672e+02	3.597e+01	4.648	9.54e-06 ***
## Condition5	1.787e+02	3.531e+01	5.060	1.73e-06 ***
## Condition6	1.820e+02	3.475e+01	5.237	8.11e-07 ***
## Condition7	1.956e+02	3.559e+01	5.496	2.62e-07 ***
## Condition8	2.063e+02	3.611e+01	5.712	9.95e-08 ***
## Condition9	2.394e+02	4.139e+01	5.784	7.20e-08 ***
## YearBuilt	1.009e+00	5.681e-01	1.775	0.078652 .
## BasementFinSF	1.211e-02	1.697e-02	0.714	0.476768
## ExteriorQFa	-1.006e+02	5.207e+01	-1.931	0.056096 .
## ExteriorQGd	-3.216e+01	2.130e+01	-1.510	0.133973
## ExteriorQTA	-3.232e+01	2.239e+01	-1.443	0.151900
## HouseStyle1Story	4.551e+00	1.143e+01	0.398	0.691239
## HouseStyle2.5Unf	-1.797e+00	2.444e+01	-0.074	0.941530
## HouseStyle2Story	-1.110e+01	9.119e+00	-1.217	0.226263
## HouseStyleSFoyer	-5.681e+00	1.414e+01	-0.402	0.688709
## HouseStyleSLvl	-3.976e+00	1.389e+01	-0.286	0.775163
## GarageSF	-2.664e-02	7.833e-02	-0.340	0.734419
## FoundationCBlock	4.698e+00	1.008e+01	0.466	0.642071
## FoundationPConc	1.302e+01	1.074e+01	1.213	0.227961
## FoundationSlab	1.687e+01	2.941e+01	0.574	0.567498
## FoundationStone	2.450e+01	2.806e+01	0.873	0.384471
## FoundationWood	-4.161e+01	2.966e+01	-1.403	0.163481
## LotArea	1.681e-04	1.836e-04	0.915	0.362050
## LotConfigCulDSac	7.015e+00	9.241e+00	0.759	0.449452
## LotConfigFR2	-2.541e+01	1.139e+01	-2.231	0.027745 *
## LotConfigFR3	-2.713e+01	3.240e+01	-0.837	0.404311
## LotConfigInside	-4.574e-01	5.062e+00	-0.090	0.928169
## GarageTypeAttchd	-2.888e+01	3.331e+01	-0.867	0.387897
## GarageTypeBasement	-2.999e+01	3.840e+01	-0.781	0.436457
## GarageTypeBuiltIn	-2.774e+01	3.479e+01	-0.798	0.426872
## GarageTypeCarPort	-7.833e+01	3.884e+01	-2.017	0.046218 *
## GarageTypeDetchd	-2.258e+01	3.166e+01	-0.713	0.477369
## GarageTypeNone	-6.769e+01	5.545e+01	-1.221	0.224847
## Bedroom	-2.395e+01	1.214e+01	-1.972	0.051110 .
## TotalRooms	1.263e+00	5.616e+00	0.225	0.822567
## ExteriorCFa	-2.703e+01	3.631e+01	-0.744	0.458204
## ExteriorCGd	-1.109e+01	2.904e+01	-0.382	0.703360
## ExteriorCTA	-8.596e-01	2.876e+01	-0.030	0.976209
## BasementHBath	-1.699e+01	1.434e+01	-1.185	0.238749
## I(GarageSF^2)	-6.440e-05	1.493e-04	-0.431	0.666981
## GarageCFa	8.627e+00	3.483e+01	0.248	0.804817
## GarageCGd	-3.961e+01	4.208e+01	-0.941	0.348680
## GarageCNone	NA	NA	NA	NA
## GarageCTA	7.820e+00	3.140e+01	0.249	0.803827
## HeatingQCFa	8.185e+00	1.460e+01	0.561	0.576229
## HeatingQCGd	1.279e+01	6.303e+00	2.029	0.044966 *
## HeatingQCTA	4.616e+00	6.657e+00	0.693	0.489569
## BasementSF	3.224e-02	2.461e-02	1.310	0.192923



```

## BasementCFa      -1.187e+01  1.411e+01 -0.841 0.402300
## BasementCGd      -7.052e+00  7.456e+00 -0.946 0.346333
## BasementCNone    -4.743e+00  3.041e+01 -0.156 0.876369
## BasementCTA              NA              NA              NA
## BasementHtEx      9.074e+00  1.244e+01  0.729 0.467506
## BasementHtFa      1.048e+01  1.654e+01  0.633 0.527864
## BasementHtGd      6.876e+00  7.472e+00  0.920 0.359471
## BasementHtNone    NA              NA              NA
## BasementHtTA              NA              NA              NA
## KitchenQFa      -3.605e+01  2.741e+01 -1.315 0.191239
## KitchenQGd      -6.624e+00  1.081e+01 -0.613 0.541368
## KitchenQTA      -8.142e+00  1.218e+01 -0.669 0.505132
## GarageQGd      -8.199e+00  2.151e+01 -0.381 0.703840
## GarageQNone              NA              NA              NA
## GarageQTA      -1.128e+01  1.375e+01 -0.820 0.413753
## ScreenPorchSF    -3.934e-02  4.675e-02 -0.842 0.401831
## Fireplaces      4.572e+00  4.099e+00  1.115 0.267191
## WoodDeckSF      3.515e-03  1.848e-02  0.190 0.849472
## FullBath      -6.706e+00  6.741e+00 -0.995 0.322085
## HalfBath      -1.375e+01  6.750e+00 -2.037 0.044100 *
## SecondSF      2.216e-01  6.379e-02  3.474 0.000739 ***
## FirstSF      2.075e-01  6.105e-02  3.398 0.000951 ***
## BasementFinALQ    -1.421e+00  8.953e+00 -0.159 0.874163
## BasementFinBLQ    -4.280e+00  9.850e+00 -0.435 0.664770
## BasementFinGLQ    -3.534e+00  8.305e+00 -0.425 0.671338
## BasementFinLwQ    -6.426e+00  1.131e+01 -0.568 0.571240
## BasementFinNone    NA              NA              NA
## BasementFinRec    -9.143e-01  1.155e+01 -0.079 0.937041
## BasementFinUnf    NA              NA              NA
## BasementUnFinSF    -1.819e-02  1.725e-02 -1.055 0.293942
## BasementFBath    -3.908e+00  5.878e+00 -0.665 0.507549
## YearRemodel      -1.235e-01  1.869e-01 -0.661 0.510279
## HeatingGasW      -1.779e+01  1.674e+01 -1.063 0.290187
## EnclosedPorchSF    -2.222e-02  3.462e-02 -0.642 0.522453
## LotFrontage      -2.930e-01  1.399e-01 -2.094 0.038602 *
## I(age_built^2)      2.779e-03  4.658e-03  0.597 0.552070
## GarageCars      -1.699e+01  2.196e+01 -0.774 0.440912
## I(GarageCars^2)    -6.113e+00  1.124e+01 -0.544 0.587617
## CentralAirY      2.857e+00  1.324e+01  0.216 0.829516
## OpenPorchSF      -3.777e-03  3.253e-02 -0.116 0.907781
## Bedroom:TotalRooms 1.649e+00  1.654e+00  0.997 0.320916
## LotArea:LotFrontage 2.398e-05  8.872e-06  2.703 0.007987 **
## GarageSF:GarageCars 6.718e-02  7.113e-02  0.944 0.347033
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 21.63 on 108 degrees of freedom
## Multiple R-squared:  0.9454, Adjusted R-squared:  0.8994
## F-statistic: 20.56 on 91 and 108 DF,  p-value: < 2.2e-16

```

```

ames10_holdout <- read.csv('AmesTest10.csv')
ames10_holdout$Quality <- factor(ames10_holdout$Quality)
ames10_holdout$Condition <- factor(ames10_holdout$Condition)

```

**Part 9: Cross-validation for your “fancy” model** We will compute the predicted Price for each of the cases in the holdout test sample, using our new model.

```
predictions <- predict(simple_model, newdata = ames10_holdout)
head(predictions)
```

```
##          1          2          3          4          5          6
## 238.1032 138.1705 185.0101 186.9975 172.7794 131.7513
```

As in part 7 we will compute the residuals for the 200 holdout cases.

```
holdout_resid <- log(ames10_holdout$Price) - predictions
```

And find the mean and standard deviation of these residuals.

```
mean(holdout_resid)
```

```
## [1] -169.0262
```

```
sd(holdout_resid)
```


```
## [1] 65.96771
```

The mean value of our residuals is incredibly close to zero which is ideal. The standard deviation of .0766 seems very low but this is in terms of  $\log(\text{Price})$  which is much smaller than price.

We will construct a plot to determine if they are normally distributed.

```
qqnorm(holdout_resid)
qqline(holdout_resid)
```

### Normal Q–Q Plot



The residuals appear much more normally distributed than in previous models. There are potentially one or two extreme points.

Compute the correlation between the predicted values and actual prices for the holdout sample. This is known as the cross-validation correlation.

```
cor(predictions, ames10_holdout$Price)
```

```
## [1] 0.9723245
```

```
cor(predictions, log(ames10_holdout$Price))
```

```
## [1] 0.947947
```

Square the cross-validation correlation to get an  $R^2$  value.

```
cor(predictions, ames10_holdout$Price)^2
```

```
## [1] 0.9454149
```

Now subtract it from the original multiple  $R^2$  of the training sample to find the shrinkage.

```
summary(simple_model)$r.squared - cor(predictions, ames10_holdout$Price)^2
```

```
## [1] 1.110223e-16
```

The shrinkage is very small so it appears we are not over fitting our training data.

**Part 10. Final Model** Again, you may choose to make some additional adjustments to your model after considering the final residual analysis. If you do so, please explain what (and why) you did and provide the `summary()` for your new final model.

```
Price ~ Quality + GroundSF + Condition + YearBuilt +  
  BasementFinSF + ExteriorQ + HouseStyle + GarageSF + Foundation +  
  LotArea + LotConfig + GarageType + Bedroom + TotalRooms +  
  ExteriorC + BasementHBath + I(GarageSF^2) + GarageC + HeatingQC +  
  BasementSF + BasementC + BasementHt + KitchenQ + GarageQ +  
  ScreenPorchSF + Fireplaces + WoodDeckSF + FullBath + HalfBath +  
  SecondSF + FirstSF + BasementFin + BasementUnFinSF + BasementFBath +  
  YearRemodel + Heating + EnclosedPorchSF + LotFrontage + I(age_built^2) +  
  GarageCars + I(GarageCars^2) + CentralAir + OpenPorchSF +  
  Bedroom:TotalRooms + LotArea:LotFrontage + GarageSF:GarageCars,  
data = NewAmesTrain_wfactor
```

Suppose that you are interested in a house in Ames that has the characteristics listed below. Construct a 95% confidence interval for the mean price of such houses.

A 2 story 11 room home, built in 1987 and remodeled in 1999 on a 21540 sq. ft. lot with 328 feet of road frontage. Overall quality is good (7) and condition is average (5). The quality and condition of the exterior are both good (Gd) and it has a poured concrete foundation. There is an 757 sq. foot basement that has excellent height, but is completely unfinished and has no bath facilities. Heating comes from a gas air furnace that is in excellent condition and there is central air conditioning. The house has 2432 sq. ft. of living space above ground, 1485 on the first floor and 947 on the second, with 4 bedrooms, 2 full and one half baths, and 1 fireplace. The 2 car, built-in garage has 588 sq. ft. of space and is average (TA) for both quality and construction. The only porches or decks is a 205 sq. ft. open porch in the front.

```
new_house_ex <- data.frame(  
  Order = 0,  
  Price = 0,  
  LotFrontage = 328,  
  LotArea = 21540,  
  LotConfig = NA,  
  HouseStyle = "2Story",  
  Quality = 7,  
  Condition = 5,  
  YearBuilt = 1987,
```

```

YearRemodel = 1999,
age_built = 2010 - 1987,
age_remodel = 2010 - 1999,
ExteriorQ = "Gd",
ExteriorC = "Gd",
Foundation = "PConc",
BasementHt = "Ex",
BasementC = "None",
BasementFin = "Unf",
BasementFinSF = 0,
BasementUnFinSF = 757,
BasementSF = 757,
Heating = "GasA",
HeatingQC = "Ex",
CentralAir = "Y",
FirstSF = 1485,
SecondSF = 947,
GroundSF = 2432,
BasementFBath = 0,
BasementHBath = 0,
FullBath = 2,
HalfBath = 1,
Bedroom = 4,
KitchenQ = NA,
TotalRooms = 11,
Fireplaces = 1,
GarageType = "BuiltIn",
GarageCars = 2,
GarageSF = 588,
GarageQ = "TA",
GarageC = "TA",
WoodDeckSF = 0,
OpenPorchSF = 205,
EnclosedPorchSF = 0,
ScreenPorchSF = 0
)

house_ex <- data.frame(new_house_ex, stringsAsFactors = TRUE)

predict.lm(simple_model, new_data = house_ex, interval = 'confidence', level = 0.95)

```

```

##          fit      lwr      upr
## 1  238.10315 215.00708 261.19923
## 2  138.17050 108.84123 167.49977
## 3  185.01011 169.69790 200.32232
## 4  186.99751 156.51768 217.47733
## 5  172.77941 149.18516 196.37365
## 6  131.75126 105.46419 158.03832
## 7  267.92839 245.93027 289.92651
## 8  185.78926 164.16522 207.41330
## 9  136.46716 116.12777 156.80654
## 10  98.30000  55.42534 141.17466
## 11 125.23664  90.40281 160.07047
## 12 147.35103 120.39318 174.30887
## 13 178.30778 149.65838 206.95718

```

## 14	245.89066	224.05794	267.72338
## 15	161.74613	139.86511	183.62715
## 16	289.64840	264.73465	314.56214
## 17	171.67475	148.68351	194.66599
## 18	288.66402	254.44704	322.88099
## 19	176.44757	144.99253	207.90262
## 20	175.19712	144.61307	205.78117
## 21	127.00000	84.12534	169.87466
## 22	290.52981	272.64251	308.41711
## 23	123.12901	101.79972	144.45831
## 24	99.63442	71.19392	128.07492
## 25	174.50203	144.47845	204.52561
## 26	237.53508	212.37463	262.69553
## 27	143.45083	112.81180	174.08986
## 28	221.04683	184.95011	257.14354
## 29	143.71220	114.51210	172.91230
## 30	262.79920	242.66594	282.93246
## 31	115.79434	85.23153	146.35714
## 32	231.88143	214.77944	248.98342
## 33	174.97169	157.13611	192.80727
## 34	219.50773	198.95500	240.06046
## 35	181.64714	152.29397	211.00032
## 36	170.73879	153.92887	187.54872
## 37	279.60950	258.11133	301.10768
## 38	255.23001	224.24678	286.21324
## 39	238.59140	211.73016	265.45264
## 40	82.70727	47.18729	118.22724
## 41	200.55850	163.00418	238.11282
## 42	137.33961	97.12178	177.55743
## 43	122.29682	101.75646	142.83719
## 44	120.92081	94.22509	147.61652
## 45	93.83956	57.86959	129.80952
## 46	113.20946	79.97635	146.44258
## 47	170.66772	149.01395	192.32149
## 48	237.18055	194.70398	279.65712
## 49	164.11061	136.00500	192.21622
## 50	155.50000	112.62534	198.37466
## 51	161.32528	133.01085	189.63972
## 52	138.68579	105.85408	171.51750
## 53	211.88021	188.72945	235.03097
## 54	201.89697	180.50738	223.28656
## 55	78.34235	41.03928	115.64541
## 56	89.58576	61.85766	117.31385
## 57	506.54447	469.31666	543.77227
## 58	139.71838	111.51721	167.91954
## 59	165.38898	136.59389	194.18406
## 60	174.21000	148.15985	200.26016
## 61	159.40273	133.36888	185.43657
## 62	242.30761	210.07120	274.54402
## 63	137.04318	107.43360	166.65276
## 64	159.93771	122.42612	197.44929
## 65	113.01572	84.40633	141.62510
## 66	115.71523	79.58168	151.84878
## 67	260.91020	240.78844	281.03195
## 68	127.45100	90.93414	163.96786

## 69	79.00000	36.12534	121.87466
## 70	208.80509	190.21256	227.39762
## 71	359.04044	328.72297	389.35791
## 72	125.10054	101.41702	148.78406
## 73	339.16003	305.74476	372.57530
## 74	162.30618	128.03218	196.58018
## 75	90.94697	57.78198	124.11197
## 76	88.10635	61.39890	114.81379
## 77	216.07663	197.80619	234.34707
## 78	173.58959	147.90190	199.27728
## 79	155.96449	126.49987	185.42912
## 80	128.46393	98.96381	157.96406
## 81	57.02896	19.92681	94.13110
## 82	165.87406	132.47397	199.27414
## 83	207.74932	192.85578	222.64287
## 84	143.82068	114.55738	173.08399
## 85	126.09480	88.24414	163.94545
## 86	146.87714	118.14440	175.60989
## 87	132.11582	106.44094	157.79070
## 88	132.48303	105.81233	159.15373
## 89	307.12254	267.28419	346.96088
## 90	232.66774	197.84586	267.48961
## 91	186.11105	164.09492	208.12718
## 92	143.00000	100.12534	185.87466
## 93	115.95621	90.85780	141.05463
## 94	191.39027	168.84989	213.93065
## 95	150.00000	107.12534	192.87466
## 96	287.29898	264.29597	310.30200
## 97	170.63702	152.32823	188.94580
## 98	173.82363	149.01786	198.62940
## 99	69.03653	32.40560	105.66746
## 100	175.30520	137.45455	213.15586
## 101	140.39733	116.37178	164.42287
## 102	58.69935	22.89808	94.50061
## 103	124.97266	89.13526	160.81006
## 104	135.23935	110.64270	159.83601
## 105	143.50310	116.04000	170.96620
## 106	111.82342	82.71786	140.92897
## 107	176.06885	154.67186	197.46584
## 108	115.89032	84.09813	147.68251
## 109	221.06229	183.55071	258.57388
## 110	118.00000	75.12534	160.87466
## 111	121.00434	96.58282	145.42586
## 112	274.47618	254.59886	294.35350
## 113	280.92580	248.19324	313.65836
## 114	149.47338	126.58167	172.36509
## 115	114.24810	86.11735	142.37884
## 116	210.64391	183.03491	238.25291
## 117	202.01792	180.48975	223.54610
## 118	150.00000	107.12534	192.87466
## 119	182.92783	147.20250	218.65316
## 120	229.50095	210.28857	248.71332
## 121	145.91211	112.64748	179.17675
## 122	112.24649	89.71910	134.77388
## 123	86.27907	47.78298	124.77517

```

## 124 204.02500 182.34187 225.70813
## 125 124.36529 94.29659 154.43400
## 126 122.97922 96.93491 149.02353
## 127 302.33863 285.31049 319.36676
## 128 264.02608 229.10699 298.94517
## 129 142.87081 114.80173 170.93989
## 130 210.43468 192.31363 228.55574
## 131 176.13338 152.14195 200.12481
## 132 183.02615 148.74727 217.30503
## 133 121.78383 100.01661 143.55105
## 134 133.03600 107.84304 158.22896
## 135 108.70928 81.82543 135.59313
## 136 141.16601 122.47950 159.85252
## 137 144.47702 112.92526 176.02879
## 138 147.95485 123.12325 172.78646
## 139 198.07295 173.54576 222.60014
## 140 132.71210 105.63879 159.78540
## 141 148.78594 118.24532 179.32655
## 142 107.66039 67.44257 147.87822
## 143 118.89412 91.03051 146.75774
## 144 161.65908 137.06946 186.24869
## 145 267.39340 236.58881 298.19799
## 146 304.16678 281.84029 326.49327
## 147 120.35575 88.05964 152.65185
## 148 193.69526 174.61211 212.77842
## 149 232.46084 208.26068 256.66099
## 150 162.93672 141.37113 184.50230
## 151 106.51556 76.76622 136.26490
## 152 205.67280 165.29701 246.04859
## 153 265.59900 232.50168 298.69632
## 154 187.28477 151.15122 223.41832
## 155 256.07065 239.47614 272.66517
## 156 103.98360 79.36134 128.60586
## 157 153.08467 119.80161 186.36773
## 158 115.18937 81.93441 148.44432
## 159 141.73564 111.94988 171.52139
## 160 126.95715 98.60502 155.30928
## 161 87.23441 49.94263 124.52618
## 162 129.37746 89.53912 169.21581
## 163 200.14524 178.51683 221.77365
## 164 114.47924 79.49701 149.46146
## 165 202.39671 185.51189 219.28153
## 166 121.03298 94.36598 147.69998
## 167 102.35584 68.15629 136.55539
## 168 153.32743 126.91272 179.74213
## 169 138.29502 105.34546 171.24459
## 170 81.00000 38.12534 123.87466
## 171 327.10850 294.22044 359.99657
## 172 102.52387 72.20970 132.83804
## 173 111.47386 77.49648 145.45124
## 174 122.58171 100.37452 144.78889
## 175 320.98777 295.41440 346.56114
## 176 113.24099 90.37787 136.10412
## 177 122.86582 95.27722 150.45442
## 178 194.78582 170.76077 218.81087

```

## 179 276.15107 244.83255 307.46958  
## 180 191.17408 174.86039 207.48777  
## 181 181.13843 163.37857 198.89829  
## 182 252.18771 222.21042 282.16501  
## 183 205.35158 190.52677 220.17640  
## 184 170.73296 142.96992 198.49599  
## 185 241.66030 214.43150 268.88909  
## 186 255.65746 229.24799 282.06694  
## 187 354.01984 328.21049 379.82918  
## 188 182.03595 156.60388 207.46802  
## 189 150.36668 131.44000 169.29337  
## 190 82.10568 48.92358 115.28778  
## 191 221.72897 202.48049 240.97744  
## 192 173.18291 151.65725 194.70858  
## 193 153.72387 122.74630 184.70144  
## 194 176.08603 148.19146 203.98060  
## 195 226.09690 201.77304 250.42077  
## 196 146.09421 116.27690 175.91153  
## 197 132.49186 107.27921 157.70451  
## 198 305.05227 280.20864 329.89591  
## 199 221.77385 199.70072 243.84697  
## 200 220.52526 194.87758 246.17294